

V5b

AES-GCM

AUTHENTICATED ENCRYPTION

CRYPTO 101: Building Blocks

©Alfred Menezes

cryptography101.ca

Overview

NIST Special Publication 800-38D
November, 2007



Recommendation for Block Cipher Modes of Operation: Galois/Counter Mode (GCM) and GMAC

Morris Dworkin

- ♦ **AES-GCM** is an authenticated encryption scheme designed by David McGrew and John Viega in 2004.
- ♦ Adopted as a NIST standard (SP 800-38D) in 2007.
- ♦ Uses the **CTR** mode of encryption and **GMAC**, a custom-designed MAC scheme.



CTR: CounTeR mode of encryption

Let $k \in_R \{0,1\}^{128}$ be the secret key shared by Alice and Bob. Let $M = (M_1, M_2, \dots, M_u)$ be a plaintext message, where each M_i is a 128-bit block and $u \leq 2^{32} - 2$.

To **encrypt** M , Alice does the following:

1. Select a **nonce** $IV \in \{0,1\}^{96}$.
2. Let $J_0 = IV \parallel 0^{31} \parallel 1$.
3. For i from 1 to u do:
 $J_i \leftarrow J_{i-1} + 1$ and compute
 $C_i = \text{AES}_k(J_i) \oplus M_i$.
4. Send $(IV, C_1, C_2, \dots, C_u)$ to Bob.


To **decrypt**, Bob does the following:

1. Let $J_0 = IV \parallel 0^{31} \parallel 1$.
2. For i from 1 to u do:
 $J_i \leftarrow J_{i-1} + 1$ and compute
 $M_i = \text{AES}_k(J_i) \oplus C_i$.

Notes on CTR mode

1. CTR mode of encryption can be viewed as a stream cipher.
2. As was the case with CBC encryption, identical plaintexts with different IVs result in different ciphertexts.
3. It is critical that the **IV should not be repeated**, but this can be difficult to achieve in practice.
4. Unlike CBC encryption, CTR encryption is **parallelizable**.
5. Note that AES^{-1} is not used.
6. The secret key can have bitlength 128, 192 or 256.

Multiplying blocks



Évariste Galois

- ♦ Let $a = a_0a_1a_2\dots a_{127}$ be a 128-bit block.
We associate the binary polynomial $a(x) = a_0 + a_1x + a_2x^2 + \dots + a_{127}x^{127} \in \mathbb{Z}_2[x]$ with a .
- ♦ Let $f(x) = 1 + x + x^2 + x^7 + x^{128}$.
- ♦ If a and b are 128-bit blocks, then define $c = a \bullet b$ to be the block corresponding to the polynomial $c(x) = a(x) \cdot b(x) \bmod f(x)$.
 - ♦ That is, $c(x)$ is the remainder upon dividing $a(x) \cdot b(x)$ by $f(x)$ in $\mathbb{Z}_2[x]$.
 - ♦ This is multiplication in the **Galois field** $GF(2^{128})$.

Galois Message Authentication Code (GMAC)

- ♦ Let $A = (A_1, A_2, \dots, A_v)$, where each A_i is a 128-bit block.
 - ♦ Let L be the bitlength of A (encoded as a 128-bit block).
 - ♦ Let $k \in_R \{0,1\}^{128}$ be the secret key.
1. Let $J_0 = IV || 0^{31} || 1$, where $IV \in \{0,1\}^{96}$ is a **nonce**.
 2. Compute $H = \text{AES}_k(0^{128})$.
 3. Let $f_A(x) = A_1x^{v+1} + A_2x^v + \dots + A_{v-1}x^3 + A_vx^2 + Lx \in GF(2^{128})[x]$.
 4. Compute the **authentication tag** $t = \text{AES}_k(J_0) \oplus f_A(H)$.
 5. Send (IV, A, t) .

Computing $f_A(H)$ using Horner's rule

- ♦ **Example:** Let $A = (A_1, A_2, A_3)$.
 - ♦ Then $f_A(x) = A_1x^4 + A_2x^3 + A_3x^2 + Lx$.
 - ♦ Hence, $f_A(H) = A_1H^4 + A_2H^3 + A_3H^2 + LH$.
 - ♦ $f_A(H)$ can be computed using **Horner's rule**:
$$f_A(H) = ((((((A_1 \cdot H) + A_2) \cdot H) + A_3) \cdot H) + L) \cdot H.$$
 - ♦ This requires three additions and four multiplications in $GF(2^{128})$.
- ♦ In general, if A has blocklength v , then computing $f_A(H)$ using Horner's rule requires v additions and $v + 1$ multiplications in $GF(2^{128})$.

Security argument

- ♦ Consider the **simplified tag**: $t' = f_A(H)$.
 - ♦ An adversary can guess the tag t' of a message A with success probability $\frac{1}{2^{128}}$.
 - ♦ She can also guess the tag t' by making a guess H' for H and computing $f_A(H')$. Her success probability is at most $\frac{v+1}{2^{128}}$, where v is the blocklength of A .
 - ♦ However, if the adversary sees a single valid message-tag pair (A, t') , she can solve the polynomial equation $f_A(H) = t'$ for H .
- ♦ To circumvent the aforementioned attack, a second secret $\text{AES}_k(J_0)$ is used to hide t' : $t = \text{AES}_k(J_0) \oplus f_A(H)$. The secret $\text{AES}_k(J_0)$ serves as a **one-time pad** for t' .

Authenticated encryption: AES-GCM

Input:

- ♦ **AAD (Additional Authenticated Data)**, also called **encryption context**: Data to be authenticated (but not encrypted): $A = (A_1, A_2, \dots, A_v)$.
- ♦ Data to be encrypted and authenticated: $M = (M_1, M_2, \dots, M_u)$, $u \leq 2^{32} - 2$.
- ♦ Secret key: $k \in_R \{0,1\}^{128}$, shared between Alice and Bob

Output: (IV, A, C, t) , where

- ♦ IV is a 96-bit initialization vector.
- ♦ $A = (A_1, A_2, \dots, A_v)$ is the additional authenticated data.
- ♦ $C = (C_1, C_2, \dots, C_u)$ is the encrypted / authenticated data.
- ♦ t is a 128-bit authentication tag.

AES-GCM encryption/authentication

Alice does the following:

1. Let $L = L_A || L_M$, where L_A, L_M are the bitlengths of A, M expressed as 64-bit integers. (L is the **length block**.)
2. Select a **nonce** $IV \in \{0,1\}^{96}$ and let $J_0 = IV || 0^{31} || 1$.
3. **Encryption**: For i from 1 to u do:
Compute $J_i = J_{i-1} + 1$ and $C_i = \text{AES}_k(J_i) \oplus M_i$.
4. **Authentication**: Compute $H = \text{AES}_k(0^{128})$.
Compute $t = \text{AES}_k(J_0) \oplus f_{A,C}(H)$.
5. **Output**: (IV, A, C, t) .

Note: $f_{A,C}(x) = A_1 x^{u+v+1} + A_2 x^{u+v} + \dots + A_{v-1} x^{u+3} + A_v x^{u+2} + C_1 x^{u+1} + C_2 x^u + \dots + C_{u-1} x^3 + C_u x^2 + Lx$

AES-GCM decryption/authentication

Upon receiving (IV, A, C, t) , Bob does the following:

1. Let $L = L_A || L_C$, where L_A, L_C are the bitlengths of A, C expressed as 64-bit integers.
2. **Authentication:** Compute $H = \text{AES}_k(0^{128})$.
Compute $t' = \text{AES}_k(J_0) \oplus f_{A,C}(H)$.
If $t' = t$ then proceed to decryption; if $t' \neq t$ then reject.
3. **Decryption:** Let $J_0 = IV || 0^{31} || 1$.
For i from 1 to u do:
Compute $J_i = J_{i-1} + 1$ and $M_i = \text{AES}_k(J_i) \oplus C_i$.
4. **Output:** (A, M) .

Some features of AES-GCM

1. Performs both authentication and encryption.
2. Supports **authentication only** (by using empty M).
3. Very fast implementations on Intel and AMD processors because of special **AES-NI** and **PCLMUL** instructions for the AES and \bullet operations.
4. Encryption and decryption can be **parallelized**.
5. AES-GCM can be used in **streaming mode**.
6. The secret key can have bitlength 128, 192 or 256.
7. Security is justified by a security proof:
 - ✦ The original McGrew-Viega security proof (2004) was wrong.
 - ✦ The proof was fixed in 2012 by Iwata-Ohashi-Minematsu.

Performance

Speed benchmarks[†] from 2018
on an Intel Xeon CPU (E3-1220 V2)
at 3.10 GHz in 64-bit mode.

[†]Relative speeds will likely be very different
on other processors.

Source: www.bearssl.org/speed.html

| Algorithm | block length | key length | digest length (bits) | speed (Mbytes/ |
|-----------------|-----------------|---------------|-------------------------|-------------------|
| ChaCha20 | — | 256 | — | 323 |
| Triple-DES | 64 | 168 | — | 21 |
| AES-128 | 128 | 128 | — | 170 |
| AES-128-NI | 128 | 128 | — | 2426 |
| AES-256 | 128 | 256 | — | 129 |
| AES-256-NI | 128 | 256 | — | 1830 |
| GMAC | 128 | 128 | 128 | 247 |
| GMAC- PCLMUL | 128 | 128 | 128 | 1741 |

IV's should not be repeated

IV's should not be repeated (with the same key k).

- ✦ Suppose an IV is reused, and an eavesdropper captures two transmissions: (IV, A_1, C_1, t_1) , (IV, A_2, C_2, t_2) . Suppose also that M_1 and M_2 have the same blocklengths, and that the eavesdropper knows M_1 .
- ✦ Then $t_1 = \text{AES}_k(J_0) \oplus f_{A_1, C_1}(H)$ and $t_2 = \text{AES}_k(J_0) \oplus f_{A_2, C_2}(H)$, so $t_1 \oplus t_2 = f_{A_1, C_1}(H) \oplus f_{A_2, C_2}(H)$.
- ✦ This polynomial equation can be quickly solved for H , and then $\text{AES}_k(J_0) = t_1 \oplus f_{A_1, C_1}(H)$ can be computed.
- ✦ Thereafter, the adversary can properly encrypt / authenticate any plaintext (of blocklength at most that of M_1).